

Least-Squares Regression

In least-squares regression, the training data contains n different training pairs $(\overline{X}_1, y_1) \dots (\overline{X}_n, y_n)$, where each \overline{X}_i is a d -dimensional representation of the data points, and each y_i is a real-valued target. The fact that the target is *real-valued* is important, because the underlying problem is then referred to as *regression* rather than *classification*. In fact, as we will see later, one can also use least-squares regression on binary targets by “pretending” that these targets are real-valued. The resulting approach is equivalent to the Widrow-Hoff learning algorithm, which is famous in the neural network literature as the second learning algorithm proposed after the perceptron.

In least-squares regression, the target variable is related to the feature variables using the relationship $\hat{y}_i = \overline{W} \cdot \overline{X}_i$.

The portion of the loss that is specific to the i th training instance is given by the following $L_i = e_i^2 = (y_i - \hat{y}_i)^2$.

The stochastic gradient-descent steps are determined by computing the gradient of e_i^2 with respect to \overline{W} , when the training pair (\overline{X}_i, y_i) is presented to the neural network.

- This gradient can be computed as follows $\frac{\partial L_i}{\partial \overline{W}} = -e_i \overline{X}_i$.

$$\Rightarrow L_i = y_i^2 + \hat{y}_i^2 - 2y_i\hat{y}_i = y_i^2 + (\overline{W} \cdot \overline{X}_i)^2 - 2y_i(\overline{W} \cdot \overline{X}_i) \Rightarrow$$

$$\Rightarrow \frac{\partial L_i}{\partial \overline{W}} = 2(\overline{W} \cdot \overline{X}_i) \cdot \overline{X}_i - 2y_i\overline{X}_i = 2(\hat{y}_i - y_i)\overline{X}_i = -2e_i\overline{X}_i$$

- Therefore, the gradient-descent updates for \overline{W} are computed using the above gradient and step-size α (hyper-parameter), $\overline{W} \Leftarrow \overline{W} + \alpha e_i \overline{X}_i = \overline{W} + \alpha(y_i - \hat{y}_i)\overline{X}_i$.

$$\Rightarrow \overline{W} \Leftarrow \overline{W} + 2\alpha e_i \overline{X}_i \text{ so we could not take into account the 2 if we think that } \alpha \Leftarrow 2\alpha.$$

With regularization, the update is as follows $\overline{W} \Leftarrow \overline{W}(1 - \alpha\lambda) + \alpha(y_i - \hat{y}_i)\overline{X}_i$, where $\lambda > 0$ is the regularization parameter.

What if we applied least-squares regression directly to minimize the squared distance of the real-valued prediction \hat{y}_i from the observed binary targets $y_i \in \{-1, +1\}$? The direct application of least-squares regression to binary targets is referred to as least-squares classification. The gradient-descent is the same as the one shown above.

This direct application of least-squares regression to binary targets is referred to as Widrow-Hoff learning.

Widrow-Hoff Learning

The loss function of the Widrow-Hoff method can be rewritten slightly from least-squares regression because of its binary responses, when working with binary responses in $\{-1, +1\}$

$$\Rightarrow L_i = (y_i - \hat{y}_i)^2 = y_i^2(y_i - \hat{y}_i)^2 = (y_i^2 - y_i\hat{y}_i)^2 = (1 - y_i\hat{y}_i)^2$$

One of the flows of this method is that it penalizes over-performance, and other methods can be shown to be closely related the Widrow-Hoff loss function by using different ways of repairing the loss so that over-performance is not penalized.

The gradient-descent updates of least-squares regression can be rewritten slightly for Widrow-Hoff learning because of binary response variables:

$$\overline{W} \Leftarrow \overline{W}(1 - \alpha \cdot \lambda) + \alpha(y_i - \hat{y}_i)\overline{X}_i = \overline{W}(1 - \alpha \cdot \lambda) + \alpha y_i(1 - y_i\hat{y}_i)\overline{X}_i$$

Closed Form Solutions

The special case of least-squares regression and classification is solvable in closed form (without gradient-descent) by using the pseudo-inverse of the $n \times d$ training data matrix D , whose rows are $\overline{X_1}, \dots, \overline{X_n}$. Let the n -dimensional column vector of dependent variables be denoted by $\overline{y} = [y_1, \dots, y_n]^T$.

- The pseudo-inverse of matrix D is defined as $D^+ = (D^T D)^{-1} D^T$.
- Then, the row-vector \overline{W} is defined by $\overline{W}^T = D^+ \overline{y}$.
- If regularization is incorporated, $\overline{W}^T = (D^T D + \lambda I)^{-1} D^T \overline{y}$.

One rarely inverts large matrices like $D^T D$. In fact, the Widrow-Hoff updates provide a very efficient way of solving the problem without using the closed form solution.